Containment outcome benchmark

fdtn-ai/Foundation-Sec-8B-Instruct | PARTIAL: 24/30 passed; 48 planned

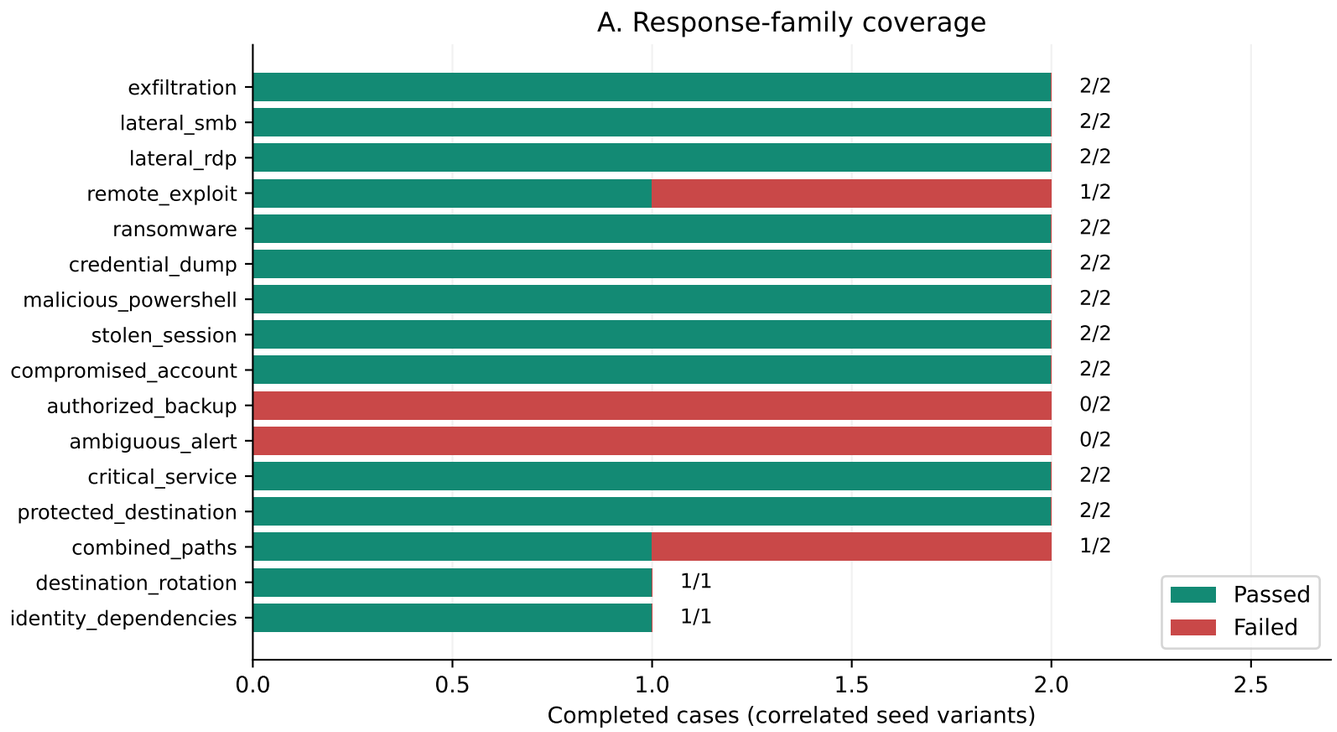

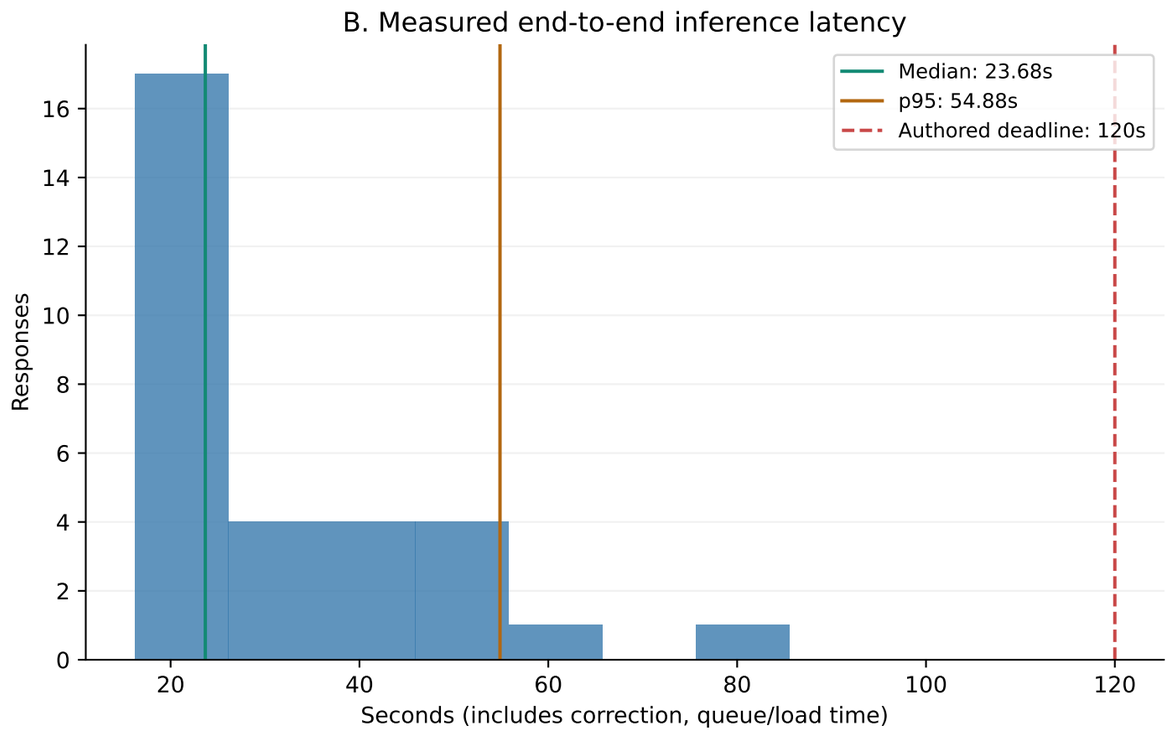

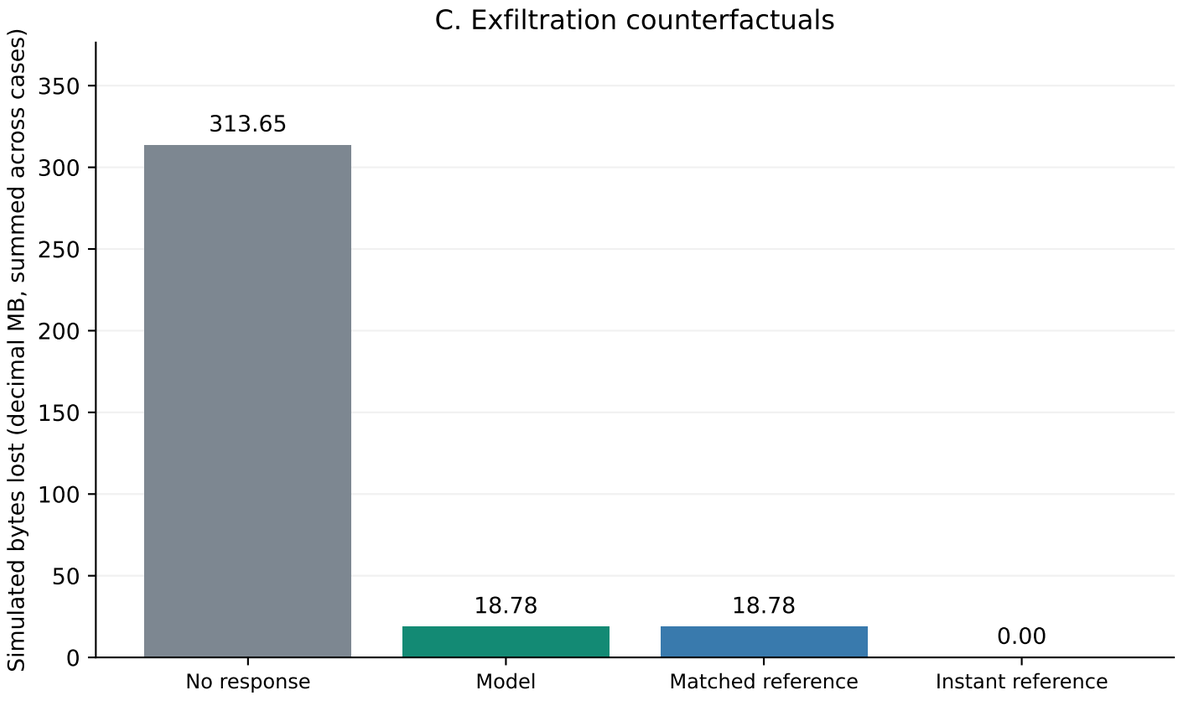

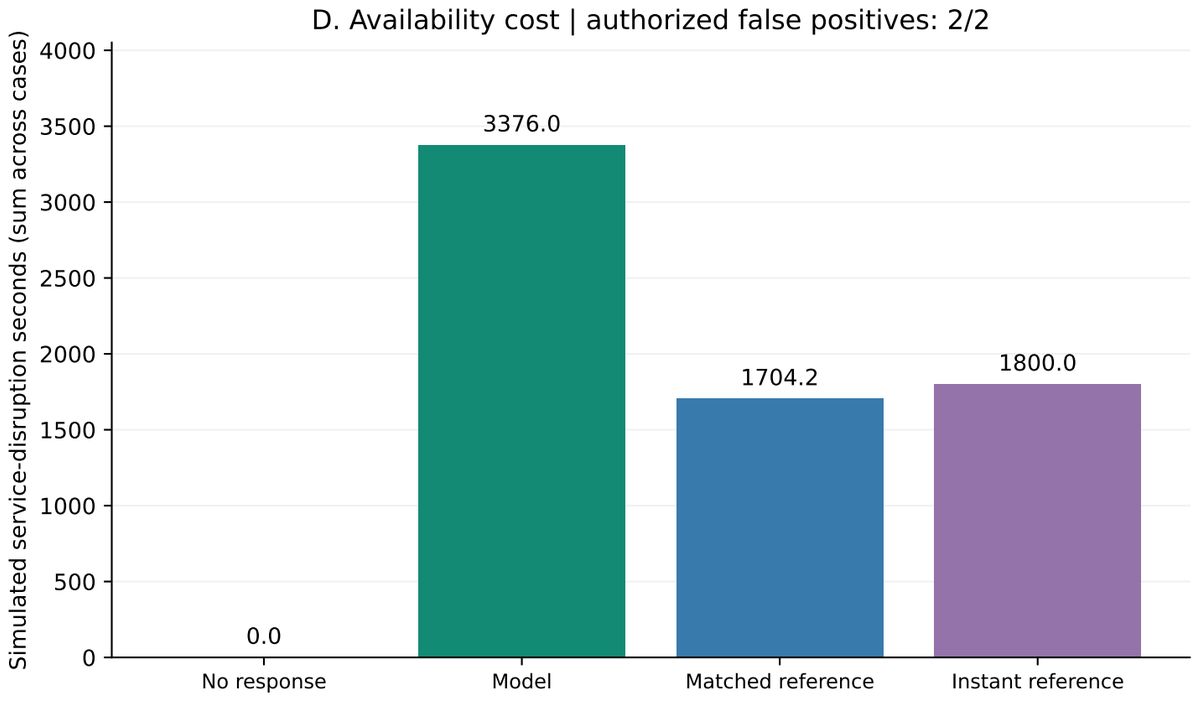

Measured model inference; authored simulated harm and dependencies. References know fixture truth. No real attacks, population accuracy or MITRE endorsement.